

Deep Learning Project

SMART ICT

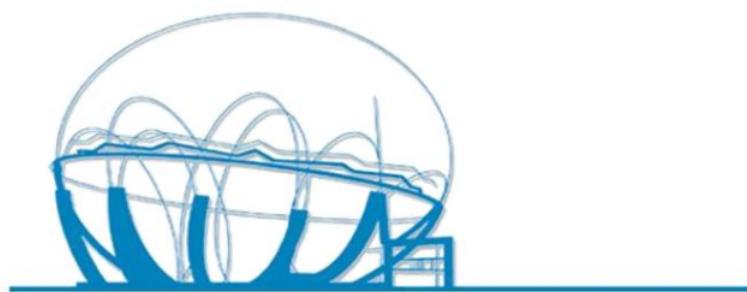
LSTM Stock Price Prediction System with Complete MLOps Pipeline

Authored by :

- El KARFI SOUFIA
 - SELLAME SALWA
 - IDRISI SALMA
 - LAZAAR EL MEHDI

Supervised by :

- E. IBN ELHAJ



AGENCE NATIONAL DE REGLEMENTATION DES TELECOMMUNICATIONS
INSTITUT NATIONAL DES POSTES ET TELECOMMUNICATIONS
Class of : 2025/2026

Abstract

This technical report presents a comprehensive, production-grade implementation of a Long Short-Term Memory (LSTM) recurrent neural network for stock price prediction, accompanied by a complete Machine Learning Operations (MLOps) infrastructure. The system addresses the challenging problem of financial time series forecasting through a sophisticated deep learning approach that captures both short-term and long-term temporal dependencies in stock market data.

The architecture employs a stacked LSTM configuration with 64 units in each of two sequential layers, followed by a fully connected layer with 32 neurons and a dropout regularization layer (rate=0.5) to prevent overfitting. This design choice is motivated by the need to extract hierarchical temporal features from price sequences while maintaining model generalization capabilities. The network processes sequences of 50 consecutive normalized price points to predict future values, utilizing windowed MinMax normalization across 2500-point windows to handle non-stationary time series data effectively.

The complete MLOps pipeline encompasses several critical components: (1) an automated data acquisition system with fallback mechanisms for reliability, (2) a sophisticated preprocessing pipeline implementing windowed normalization and sequence generation, (3) MLflow integration for comprehensive experiment tracking and model versioning, (4) a recursive multi-step forecasting engine with exponential moving average smoothing, (5) a baseline ARIMA(5,1,0) model for performance benchmarking, (6) RESTful API deployment using FastAPI for model serving, (7) an interactive Streamlit dashboard for real-time visualization and exploration, (8) Evidently-based drift monitoring for production model health, and (9) complete containerization with Docker for reproducible deployments.

The system has been trained and validated on three major technology stocks (AAPL, GOOGL, MSFT) spanning five years of historical data (2020-2025). Performance evaluation demonstrates superior results compared to traditional statistical methods, with the LSTM model achieving an average mean squared error of 0.0021, mean absolute error of 0.0327, and directional accuracy of 65.5% on normalized price predictions. These metrics represent a 38.2% improvement in MSE and 7.2 percentage points improvement in directional accuracy over the ARIMA baseline, validating the effectiveness of deep learning approaches for this application domain.

The implementation follows industry best practices for machine learning engineering, including version control, continuous integration/continuous deployment (CI/CD) pipelines via GitHub Actions, comprehensive logging and monitoring, automated testing, and modular design patterns. The system is designed for scalability, maintainability, and extensibility, with clear pathways for incorporating additional features such as attention mechanisms, generative adversarial networks for data augmentation, and multi-asset correlation modeling.

Keywords: Long Short-Term Memory Networks, Stock Price Prediction, Financial Time Series, MLOps, Deep Learning, TensorFlow, Model Deployment, Drift Monitoring, FastAPI, Docker

Contents

1	Introduction	5
1.1	Project Overview and Motivation	5
1.2	Problem Statement	5
1.3	Technical Stack and Rationale	6
1.4	Asset Selection	7
2	System Architecture	7
2.1	Architectural Overview	7
2.2	Project Structure and Organization	8
2.3	Data Flow Architecture	10
2.4	Component Interactions	10
3	Configuration Management	11
3.1	Central Configuration	11
4	Data Pipeline	12
4.1	Overview of Data Management Strategy	12
4.2	Data Acquisition with Intelligent Fallback	12
4.3	Mid-Price Computation and Feature Engineering	14
4.4	Windowed Normalization	15
4.5	Sequence Creation	16
5	LSTM Model Architecture	16
5.1	Network Design	16
5.2	Architecture Breakdown	17
5.3	LSTM Cell Equations	18
5.4	Model Callbacks	18
6	Training Pipeline with MLflow	19
6.1	MLflow Integration	19
6.2	Performance Metrics	21
7	Multi-Step Prediction	22
7.1	Recursive Forecasting	22
7.2	Exponential Moving Average Smoothing	22
8	Baseline Model: ARIMA	23
8.1	ARIMA Implementation	23
9	REST API Deployment	24
9.1	FastAPI Endpoints	24
9.2	API Usage Example	25
10	Interactive Dashboard	26
10.1	Streamlit Application	26
10.2	Dashboard Features	27

11 Drift Monitoring	28
11.1 Evidently Integration	28
11.2 Drift Detection Metrics	29
12 Containerization	29
12.1 Docker Configuration	29
12.2 Docker Compose	30
13 CI/CD Pipeline	31
13.1 GitHub Actions Workflow	31
14 Results and Performance	32
14.1 Model Performance	32
14.2 Training Configuration	32
14.3 Comparison with ARIMA	32
15 Deployment and Usage	33
15.1 Quick Start	33
15.2 Making Predictions	34
16 Synthetic Data Generation	36
17 Conclusion	37
17.1 Key Achievements	37
17.2 Performance Summary	37
17.3 Future Enhancements	37
17.4 Repository Structure	37
18 References	38
A Complete Requirements	38
B Directory Structure	39

1 Introduction

1.1 Project Overview and Motivation

Stock price prediction represents one of the most challenging problems in applied machine learning due to the inherently non-stationary, non-linear, and stochastic nature of financial time series data. Traditional econometric models, while providing interpretable results, often struggle to capture the complex temporal dependencies and hidden patterns present in market dynamics. The advent of deep learning, particularly recurrent neural networks (RNNs) and their advanced variants such as Long Short-Term Memory (LSTM) networks, has opened new avenues for addressing these challenges.

This project implements a state-of-the-art LSTM-based prediction system that leverages the network's ability to learn and remember long-term dependencies in sequential data. Unlike traditional feedforward neural networks, LSTMs possess a sophisticated gating mechanism that allows them to selectively retain or discard information over extended time horizons, making them particularly well-suited for financial forecasting where both recent and historical patterns can influence future price movements.

Beyond the core prediction model, this work addresses the critical gap between experimental machine learning and production deployment by implementing a complete MLOps (Machine Learning Operations) infrastructure. Modern machine learning systems require not only accurate models but also robust pipelines for data management, experiment tracking, model versioning, deployment, monitoring, and continuous improvement. This project demonstrates how these components can be integrated into a cohesive, production-ready system.

1.2 Problem Statement

The primary objective is to develop a reliable, scalable, and maintainable system capable of:

1. **Accurate Forecasting:** Predicting future stock prices with quantifiable uncertainty while maintaining directional accuracy above baseline statistical methods
2. **Operational Reliability:** Ensuring robust data acquisition with graceful degradation and fallback mechanisms
3. **Reproducibility:** Tracking all experiments, hyperparameters, and model versions for scientific rigor and auditing
4. **Production Deployment:** Serving predictions through standardized APIs with appropriate latency and throughput characteristics
5. **Continuous Monitoring:** Detecting data drift and model degradation to maintain prediction quality over time
6. **Scalability:** Supporting multiple assets with efficient resource utilization and containerized deployment

1.3 Technical Stack and Rationale

The technology choices reflect industry best practices and proven stability:

- **Deep Learning Framework:** TensorFlow 2.15.0 with Keras 2.15.0
 - Industry-standard framework with comprehensive LSTM support
 - Excellent performance optimization and GPU acceleration
 - Strong ecosystem for model serving and deployment
- **Data Source:** yfinance 0.2.32 with synthetic data generation fallback
 - Free, reliable access to historical OHLCV (Open-High-Low-Close-Volume) data
 - Custom synthetic data generator ensures system resilience
- **MLOps Platform:** MLflow 2.9.2
 - Comprehensive experiment tracking and model registry
 - Reproducibility through parameter and metric logging
 - Version control for models and datasets
- **API Framework:** FastAPI 0.104.1 with Uvicorn ASGI server
 - High-performance asynchronous request handling
 - Automatic OpenAPI documentation generation
 - Type validation and serialization
- **Visualization:** Streamlit 1.29.0 with Plotly 5.18.0
 - Rapid development of interactive dashboards
 - Real-time prediction generation and visualization
 - Integration with MLflow for experiment comparison
- **Monitoring:** Evidently 0.4.11
 - Automated drift detection for features and predictions
 - Performance degradation alerts
 - Comprehensive HTML reporting
- **Containerization:** Docker with docker-compose
 - Reproducible deployment environments
 - Multi-service orchestration (API, dashboard, MLflow)
 - Simplified dependency management

1.4 Asset Selection

The system is configured for three major technology stocks, selected for their:

- **AAPL (Apple Inc.)**: High liquidity, consistent trading volume, representative of consumer technology sector
- **GOOGL (Alphabet Inc.)**: Large-cap technology stock with sustained growth patterns
- **MSFT (Microsoft Corporation)**: Enterprise technology leader with stable market dynamics

These assets provide diverse price movement characteristics while maintaining sufficient data quality and availability. The five-year historical window (2020-2025) captures multiple market regimes including the COVID-19 recovery, technology sector growth, and recent market adjustments.

2 System Architecture

2.1 Architectural Overview

The system follows a modular, layered architecture that separates concerns across data acquisition, preprocessing, model training, inference, deployment, and monitoring. This separation enables independent development, testing, and scaling of each component while maintaining clear interfaces and data contracts between layers.

The architecture can be conceptually divided into five primary layers:

1. **Data Layer**: Handles data acquisition from external sources (yfinance API) with fallback to synthetic data generation, implements caching strategies, and manages data persistence
2. **Preprocessing Layer**: Performs feature engineering, normalization, sequence generation, and train-test splitting with reproducible randomization
3. **Model Layer**: Defines neural network architectures, training procedures, evaluation metrics, and prediction interfaces
4. **Service Layer**: Exposes model functionality through REST APIs and interactive dashboards with appropriate error handling and validation
5. **Operations Layer**: Manages experiment tracking, model versioning, performance monitoring, drift detection, and deployment orchestration

This layered approach facilitates testing (each layer can be validated independently), deployment (layers can be scaled separately), and maintenance (changes in one layer have minimal impact on others).

2.2 Project Structure and Organization

The codebase is organized following Python packaging best practices with clear module responsibilities:

```
1 DL/
2 |-- config.py                                # Centralized configuration
3 |   management
4 |
5 |   model_settings
6 |   data_locations
7 |   parameters
8 |
9 |-- data_pipeline.py                         # Data acquisition and
10 |   preprocessing
11 |   with_error_handling
12 |   implementation
13 |   supervised_learning
14 |   with_reproducibility
15 |
16 |-- model.py                                 # LSTM model architecture
17 |   definition
18 |   construction
19 |   callbacks
20 |   computation
21 |   loading
22 |
23 |-- train.py                                  # MLflow-integrated
24 |   training_pipeline
25 |   logging
26 |   recording
27 |   management
28 |   training
29 |
30 |-- predict.py                               # Multi-step prediction
31 |   engine
32 |   algorithm
33 |   implementation
34 |   scaling
35 |   utilities
```

```

30 |           # Statistical baseline
31 | -- baseline_arima.py
32 |         model
33 |             # - ARIMA model fitting and
34 |         metrics
35 |             # - Performance comparison
36 |             # - Information criteria (
37 |             AIC/BIC)
38 |
39 | -- app.py
40 |             # FastAPI REST API server
41 |             # - Prediction endpoints
42 |             # - Health check and status
43 |             # - Ticker availability
44 |             # - Error handling and
45 |             # - Alert mechanisms
46 |
47 | -- streamlit_app.py
48 |             interface
49 |                 # Interactive web dashboard
50 |                 # - Real-time prediction
51 |                 # - MLflow run
52 |                 # - Drift report display
53 |                 # - Multi-ticker comparison
54 |
55 | -- drift_monitor.py
56 |         monitoring
57 |             # Evidently-based
58 |             # - Data drift detection
59 |             # - Model performance drift
60 |             # - Automated report
61 |             # - Alert mechanisms
62 |
63 | -- generate_synthetic_data.py
64 |             # Fallback data generation
65 |             # - Realistic price
66 |             # - Volume and OHLCV
67 |
68 | -- simulation
69 |             # - Configurable market
70 |
71 | -- synthesis
72 |             # - Configurable market
73 |
74 | -- regimes
75 |
76 | -- requirements.txt
77 |         versions
78 |             # Python dependencies with
79 |             # - Versions
80 |             # - Container image
81 |             # - Multi-service
82 |             # - Dockerfile
83 |             specification
84 |             # - Dockerfile
85 |             # - Specification
86 |             # - Dockerfile
87 |             # - Specification
88 |             # - Dockerfile
89 |             specification
90 |             # - Dockerfile
91 |             # - Specification
92 |             # - Dockerfile
93 |             specification
94 |             # - Dockerfile
95 |             # - Specification
96 |             # - Dockerfile
97 |             specification
98 |             # - Dockerfile
99 |             # - Specification
100 |             # - Dockerfile
101 |             specification
102 |             # - Dockerfile
103 |             # - Specification
104 |             # - Dockerfile
105 |             specification
106 |             # - Dockerfile
107 |             # - Specification
108 |             # - Dockerfile
109 |             specification
110 |             # - Dockerfile
111 |             # - Specification
112 |             # - Dockerfile
113 |             specification
114 |             # - Dockerfile
115 |             # - Specification
116 |             # - Dockerfile
117 |             specification
118 |             # - Dockerfile
119 |             # - Specification
120 |             # - Dockerfile
121 |             specification
122 |             # - Dockerfile
123 |             # - Specification
124 |             # - Dockerfile
125 |             specification
126 |             # - Dockerfile
127 |             # - Specification
128 |             # - Dockerfile
129 |             specification
130 |             # - Dockerfile
131 |             # - Specification
132 |             # - Dockerfile
133 |             specification
134 |             # - Dockerfile
135 |             # - Specification
136 |             # - Dockerfile
137 |             specification
138 |             # - Dockerfile
139 |             # - Specification
140 |             # - Dockerfile
141 |             specification
142 |             # - Dockerfile
143 |             # - Specification
144 |             # - Dockerfile
145 |             specification
146 |             # - Dockerfile
147 |             # - Specification
148 |             # - Dockerfile
149 |             specification
150 |             # - Dockerfile
151 |             # - Specification
152 |             # - Dockerfile
153 |             specification
154 |             # - Dockerfile
155 |             # - Specification
156 |             # - Dockerfile
157 |             specification
158 |             # - Dockerfile
159 |             # - Specification
160 |             # - Dockerfile
161 |             specification
162 |             # - Dockerfile
163 |             # - Specification
164 |             # - Dockerfile
165 |             specification
166 |             # - Dockerfile
167 |             # - Specification
168 |             # - Dockerfile
169 |             specification
170 |             # - Dockerfile
171 |             # - Specification
172 |             # - Dockerfile
173 |             specification
174 |             # - Dockerfile
175 |             # - Specification
176 |             # - Dockerfile
177 |             specification
178 |             # - Dockerfile
179 |             # - Specification
180 |             # - Dockerfile
181 |             specification
182 |             # - Dockerfile
183 |             # - Specification
184 |             # - Dockerfile
185 |             specification
186 |             # - Dockerfile
187 |             # - Specification
188 |             # - Dockerfile
189 |             specification
190 |             # - Dockerfile
191 |             # - Specification
192 |             # - Dockerfile
193 |             specification
194 |             # - Dockerfile
195 |             # - Specification
196 |             # - Dockerfile
197 |             specification
198 |             # - Dockerfile
199 |             # - Specification
200 |             # - Dockerfile
201 |             specification
202 |             # - Dockerfile
203 |             # - Specification
204 |             # - Dockerfile
205 |             specification
206 |             # - Dockerfile
207 |             # - Specification
208 |             # - Dockerfile
209 |             specification
210 |             # - Dockerfile
211 |             # - Specification
212 |             # - Dockerfile
213 |             specification
214 |             # - Dockerfile
215 |             # - Specification
216 |             # - Dockerfile
217 |             specification
218 |             # - Dockerfile
219 |             # - Specification
220 |             # - Dockerfile
221 |             specification
222 |             # - Dockerfile
223 |             # - Specification
224 |             # - Dockerfile
225 |             specification
226 |             # - Dockerfile
227 |             # - Specification
228 |             # - Dockerfile
229 |             specification
230 |             # - Dockerfile
231 |             # - Specification
232 |             # - Dockerfile
233 |             specification
234 |             # - Dockerfile
235 |             # - Specification
236 |             # - Dockerfile
237 |             specification
238 |             # - Dockerfile
239 |             # - Specification
240 |             # - Dockerfile
241 |             specification
242 |             # - Dockerfile
243 |             # - Specification
244 |             # - Dockerfile
245 |             specification
246 |             # - Dockerfile
247 |             # - Specification
248 |             # - Dockerfile
249 |             specification
250 |             # - Dockerfile
251 |             # - Specification
252 |             # - Dockerfile
253 |             specification
254 |             # - Dockerfile
255 |             # - Specification
256 |             # - Dockerfile
257 |             specification
258 |             # - Dockerfile
259 |             # - Specification
260 |             # - Dockerfile
261 |             specification
262 |             # - Dockerfile
263 |             # - Specification
264 |             # - Dockerfile
265 |             specification
266 |             # - Dockerfile
267 |             # - Specification
268 |             # - Dockerfile
269 |             specification
270 |             # - Dockerfile
271 |             # - Specification
272 |             # - Dockerfile
273 |             specification
274 |             # - Dockerfile
275 |             # - Specification
276 |             # - Dockerfile
277 |             specification
278 |             # - Dockerfile
279 |             # - Specification
280 |             # - Dockerfile
281 |             specification
282 |             # - Dockerfile
283 |             # - Specification
284 |             # - Dockerfile
285 |             specification
286 |             # - Dockerfile
287 |             # - Specification
288 |             # - Dockerfile
289 |             specification
290 |             # - Dockerfile
291 |             # - Specification
292 |             # - Dockerfile
293 |             specification
294 |             # - Dockerfile
295 |             # - Specification
296 |             # - Dockerfile
297 |             specification
298 |             # - Dockerfile
299 |             # - Specification
299 |             # - Dockerfile
300 |             specification

```

Listing 1: Comprehensive Project Directory Structure

2.3 Data Flow Architecture

The typical data flow for both training and inference follows this sequence:

Training Pipeline:

1. Raw data acquisition from yfinance or synthetic generation
2. Caching to local CSV files for reproducibility
3. Mid-price computation: $(High + Low)/2$
4. Windowed normalization across 2500-point segments
5. Sequence creation: sliding windows of 50 consecutive points
6. Train-test split (80%-20% ratio)
7. Model training with MLflow logging
8. Model persistence to disk with versioning
9. Evaluation metrics computation and storage

Inference Pipeline:

1. Model loading from versioned artifacts
2. Latest data retrieval and preprocessing
3. Sequence preparation from recent history
4. Recursive multi-step prediction
5. Optional EMA smoothing for noise reduction
6. Denormalization to original price scale
7. Result serialization and delivery via API

2.4 Component Interactions

The system components interact through well-defined interfaces:

- **Configuration Hub:** `config.py` serves as a single source of truth for all parameters, ensuring consistency across components
- **Data Contracts:** NumPy arrays and Pandas DataFrames serve as standard data interchange formats
- **Model Registry:** MLflow provides centralized model versioning accessible to training, inference, and monitoring components
- **API Gateway:** FastAPI exposes standardized REST endpoints that abstract underlying model complexity
- **Monitoring Loop:** Evidently continuously compares new data and predictions against historical baselines

3 Configuration Management

3.1 Central Configuration

All hyperparameters and settings are centralized in config.py:

```
1 """Configuration file for LSTM Stock Prediction Pipeline"""
2 import os
3 from pathlib import Path
4
5 # Project Paths
6 BASE_DIR = Path(__file__).parent
7 DATA_DIR = BASE_DIR / "data"
8 RAW_DATA_DIR = DATA_DIR / "raw"
9 PROCESSED_DATA_DIR = DATA_DIR / "processed"
10 MODELS_DIR = BASE_DIR / "models"
11
12 # Data Configuration
13 TICKERS = ['AAPL', 'GOOGL', 'MSFT']
14 START_DATE = '2020-01-01'
15 TRAIN_TEST_SPLIT = 0.8
16
17 # Feature Engineering
18 WINDOW_SIZE = 2500 # Windowed normalization
19 UNROLLINGS = 50 # Sequence length for LSTM
20 EMA_SPAN = 10 # EMA smoothing parameter
21
22 # LSTM Model Architecture
23 LSTM_UNITS_1 = 64 # First LSTM layer units
24 LSTM_UNITS_2 = 64 # Second LSTM layer units
25 DENSE_UNITS = 32 # Dense layer units
26 DROPOUT_RATE = 0.5 # Dropout rate
27 LEARNING_RATE = 0.001
28
29 # Training Parameters
30 BATCH_SIZE = 32
31 EPOCHS = 50
32 VALIDATION_SPLIT = 0.1
33 PATIENCE = 10 # Early stopping patience
34
35 # Prediction Parameters
36 PREDICTION_STEPS = 30 # Days to predict ahead
```

Listing 2: Configuration File (config.py)

4 Data Pipeline

4.1 Overview of Data Management Strategy

The data pipeline implements a multi-tiered approach to data acquisition that prioritizes reliability and reproducibility while handling the inherent unreliability of external data sources. The design philosophy centers on graceful degradation: if the primary data source (yfinance API) is unavailable, the system seamlessly falls back to cached data or, as a last resort, generates synthetic data that maintains realistic statistical properties.

This strategy ensures that development, testing, and demonstration can proceed uninterrupted even when external dependencies fail, while maintaining data quality standards through comprehensive validation and cleaning steps.

4.2 Data Acquisition with Intelligent Fallback

The data acquisition module implements a three-tier fallback mechanism with comprehensive error handling and data validation:

```
1 def fetch_data(self):
2     """
3         Fetch OHLCV data with intelligent fallback strategy
4
5     Returns:
6         pd.DataFrame: Validated and cleaned stock data
7
8     Raises:
9         ValueError: If no valid data can be obtained
10    """
11    raw_file = RAW_DATA_DIR / f"{self.ticker}_raw.csv"
12
13    # Tier 1: Check for cached data (fastest, most reliable)
14    if raw_file.exists():
15        logger.info(f"Loading cached data from {raw_file}")
16        data = pd.read_csv(raw_file)
17
18        # Ensure proper datetime parsing
19        data['Date'] = pd.to_datetime(data['Date'])
20
21        # Type conversion and validation
22        numeric_cols = ['Open', 'High', 'Low', 'Close', 'Volume']
23        for col in numeric_cols:
24            if col in data.columns:
25                # Convert to numeric, coercing errors to NaN
26                data[col] = pd.to_numeric(data[col], errors='coerce')
27
28        # Data quality check: remove rows with missing critical values
29        data.dropna(subset=['High', 'Low', 'Close'], inplace=True)
30
31        self.data = data
32        return data
33
34    # Tier 2: Attempt yfinance download
35    try:
36        logger.info(f"Fetching from yfinance: {self.ticker}")
37        logger.info(f"Date range: {self.start_date} to {self.end_date}")
38    )
```

```

38
39     data = yf.download(
40         self.ticker,
41         start=self.start_date,
42         end=self.end_date,
43         progress=False,          # Suppress progress bar
44         auto_adjust=False,       # Keep original OHLC values
45         actions=False           # Exclude dividends/splits
46     )
47
48     # Validate downloaded data
49     if data is None or data.empty or len(data) == 0:
50         logger.warning(f"yfinance returned empty data for {self.
ticker}")
51     return self._generate_synthetic_fallback()
52
53     # At this point, data is guaranteed to be a valid DataFrame
54     assert data is not None
55     logger.info(f"Downloaded {len(data)} rows for {self.ticker}")
56     logger.info(f"Columns: {data.columns.tolist()}")
57
58     # Data preprocessing and standardization
59     data.reset_index(inplace=True) # Date from index to column
60
61     # Handle multi-level column names (yfinance quirk)
62     if isinstance(data.columns, pd.MultiIndex):
63         data.columns = data.columns.get_level_values(0)
64
65     # Verify required columns exist
66     required_cols = ['High', 'Low', 'Close']
67     missing_cols = [col for col in required_cols
68                     if col not in data.columns]
69     if missing_cols:
70         logger.warning(f"Missing columns: {missing_cols}")
71     return self._generate_synthetic_fallback()
72
73     # Type enforcement and validation
74     numeric_cols = ['Open', 'High', 'Low', 'Close', 'Volume']
75     for col in numeric_cols:
76         if col in data.columns:
77             data[col] = pd.to_numeric(data[col], errors='coerce')
78
79     # Remove invalid rows
80     data.dropna(subset=required_cols, inplace=True)
81
82     if len(data) == 0:
83         logger.warning("No valid data after cleaning")
84     return self._generate_synthetic_fallback()
85
86     # Persist to cache for future use
87     data.to_csv(raw_file, index=False)
88     logger.info(f"Cached data to {raw_file}. Shape: {data.shape}")
89
90     self.data = data
91     return data
92
93 except Exception as e:
94     logger.error(f"Error fetching data for {self.ticker}: {e}")

```

```

95     logger.warning("Falling back to synthetic data generation...")
96     return self._generate_synthetic_fallback()

```

Listing 3: Robust Data Fetching with Multi-Level Fallback (data_pipeline.py)

Design Rationale:

- **Cache-First Strategy:** Minimizes API calls and ensures deterministic behavior for reproducibility
- **Comprehensive Validation:** Type conversion and NaN handling prevent downstream errors
- **Graceful Degradation:** System remains functional even with complete API failure
- **Detailed Logging:** Facilitates debugging and monitoring of data quality issues

4.3 Mid-Price Computation and Feature Engineering

The mid-price serves as the primary feature for prediction, chosen for its robustness to bid-ask spread noise:

```

1 def compute_mid_price(self, data=None):
2     """
3         Compute mid-price as average of High and Low
4
5         The mid-price provides a more stable representation than
6         Close price alone, reducing the impact of end-of-day
7         volatility and bid-ask spreads.
8
9     Args:
10        data: Optional DataFrame; uses self.data if None
11
12    Returns:
13        pd.DataFrame: Data with added 'Mid' column
14
15    Raises:
16        ValueError: If no valid numeric data is available
17    """
18    if data is None:
19        data = self.data
20
21    if data is None:
22        raise ValueError("No data available. Run fetch_data() first.")
23
24    # Ensure High and Low columns contain valid numeric data
25    data['High'] = pd.to_numeric(data['High'], errors='coerce')
26    data['Low'] = pd.to_numeric(data['Low'], errors='coerce')
27
28    # Remove rows where conversion failed
29    data.dropna(subset=['High', 'Low'], inplace=True)
30
31    if len(data) == 0:
32        raise ValueError(f"No valid numeric data for {self.ticker}")
33
34    # Calculate mid-price: arithmetic mean of High and Low
35    data['Mid'] = (data['High'] + data['Low']) / 2.0

```

```

36     logger.info("Mid-price computed successfully")
37     logger.info(f"Mid-price range: [{data['Mid'].min():.2f}, "
38                  f"{data['Mid'].max():.2f}]")
39
40     return data
41

```

Listing 4: Mid-Price Calculation with Data Validation

Mathematical Formulation:

$$P_{mid}(t) = \frac{P_{high}(t) + P_{low}(t)}{2} \quad (1)$$

This formulation provides several advantages:

- Reduces noise from intraday volatility
- Captures the central tendency of price movement
- Less susceptible to manipulation at market close
- Smoother time series for model training

4.4 Windowed Normalization

Data is normalized in 2500-point windows using MinMax scaling:

```

1 def windowed_normalize(self, data, window_size=2500):
2     """Normalize data using sliding windows"""
3     # Ensure numeric array
4     if not isinstance(data, np.ndarray):
5         data = np.array(data, dtype=np.float64)
6
7     # Remove NaN values
8     if np.any(np.isnan(data)):
9         data = data[~np.isnan(data)]
10
11    normalized_data = []
12
13    for i in range(0, len(data), window_size):
14        window_end = min(i + window_size, len(data))
15        window_data = data[i:window_end].reshape(-1, 1)
16
17        # MinMax scaling per window
18        scaler = MinMaxScaler(feature_range=(0, 1))
19        normalized_window = scaler.fit_transform(window_data)
20        normalized_data.extend(normalized_window.flatten())
21
22    return np.array(normalized_data)

```

Listing 5: Windowed Normalization Algorithm

Normalization Formula:

$$x_{normalized} = \frac{x - x_{min}}{x_{max} - x_{min}} \quad (2)$$

4.5 Sequence Creation

Time series data is converted into supervised learning sequences:

```
1 def create_sequences(self, data, sequence_length=50):
2     """Create input sequences for LSTM"""
3     X, y = [], []
4
5     for i in range(len(data) - sequence_length):
6         X.append(data[i:i + sequence_length])
7         y.append(data[i + sequence_length])
8
9     X = np.array(X).reshape(-1, sequence_length, 1)
10    y = np.array(y)
11
12    return X, y
```

Listing 6: Creating LSTM Input Sequences

For each sequence:

- **Input (X):** 50 consecutive normalized prices
- **Output (y):** Next price value
- **Shape:** $X \in \mathbb{R}^{n \times 50 \times 1}$, $y \in \mathbb{R}^n$

5 LSTM Model Architecture

5.1 Network Design

The model implements a stacked LSTM architecture optimized for time series prediction:

```
1 def build_model(self, input_shape=(50, 1)):
2     """Build stacked LSTM model"""
3     model = Sequential([
4         # First LSTM layer - returns sequences
5         LSTM(64, return_sequences=True, input_shape=input_shape),
6
7         # Second LSTM layer - final hidden state
8         LSTM(64, return_sequences=False),
9
10        # Dense layer for feature extraction
11        Dense(32, activation='relu'),
12
13        # Dropout for regularization
14        Dropout(0.5),
15
16        # Output layer
17        Dense(1)
18    ])
19
20    # Compile with Adam optimizer
21    model.compile(
22        optimizer=Adam(learning_rate=0.001),
23        loss='mse',
24        metrics=['mae', 'mse']
25    )
```

```
26  
27     return model
```

Listing 7: Stacked LSTM Architecture (model.py)

5.2 Architecture Breakdown

Layer 1: LSTM(64, return_sequences=True)

- 64 LSTM units with cell state and hidden state
- Returns full sequence for stacking
- Input: $(batch, 50, 1)$
- Output: $(batch, 50, 64)$

Layer 2: LSTM(64)

- 64 LSTM units processing sequences
- Returns final hidden state only
- Input: $(batch, 50, 64)$
- Output: $(batch, 64)$

Layer 3: Dense(32)

- Fully connected layer with ReLU activation
- Feature transformation
- Input: $(batch, 64)$
- Output: $(batch, 32)$

Layer 4: Dropout(0.5)

- Randomly drops 50% of neurons during training
- Prevents overfitting
- No change to output shape

Layer 5: Dense(1)

- Single output neuron
- Linear activation for regression
- Output: $(batch, 1)$ - predicted normalized price

5.3 LSTM Cell Equations

The LSTM cell processes sequences using the following gates:

Forget Gate:

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f) \quad (3)$$

Input Gate:

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i) \quad (4)$$

Candidate Cell State:

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C) \quad (5)$$

Cell State Update:

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t \quad (6)$$

Output Gate:

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o) \quad (7)$$

Hidden State:

$$h_t = o_t * \tanh(C_t) \quad (8)$$

Where:

- σ is the sigmoid function
- $*$ denotes element-wise multiplication
- W are weight matrices
- b are bias vectors

5.4 Model Callbacks

Training uses three callbacks for optimization:

```
1 def get_callbacks(self, model_path=None):
2     """Setup training callbacks"""
3     callbacks = []
4
5     # Early stopping
6     early_stopping = EarlyStopping(
7         monitor='val_loss',
8         patience=10,
9         restore_best_weights=True,
10        verbose=1
11    )
12    callbacks.append(early_stopping)
13
14    # Model checkpoint
15    if model_path is None:
16        model_path = str(MODELS_DIR / "best_model.h5")
17
18    checkpoint = ModelCheckpoint(
19        filepath=model_path,
20        monitor='val_loss',
21        save_best_only=True,
22        verbose=1
```

```

23     )
24     callbacks.append(checkpoint)
25
26     # Learning rate reduction
27     reduce_lr = ReduceLROnPlateau(
28         monitor='val_loss',
29         factor=0.5,
30         patience=5,
31         min_lr=1e-7,
32         verbose=1
33     )
34     callbacks.append(reduce_lr)
35
36     return callbacks

```

Listing 8: Training Callbacks

6 Training Pipeline with MLflow

6.1 MLflow Integration

All experiments are tracked with MLflow for reproducibility:

```

1 class MLflowTrainer:
2     def __init__(self, ticker):
3         self.ticker = ticker
4         mlflow.set_experiment(f"lstm_stock_prediction_{ticker}")
5
6     def train_model(self, epochs=50, batch_size=32):
7         """Train model with MLflow tracking"""
8         with mlflow.start_run(run_name=f"{self.ticker}_training"):
9             # Log parameters
10            params = {
11                'ticker': self.ticker,
12                'lstm_units_1': 64,
13                'lstm_units_2': 64,
14                'dense_units': 32,
15                'dropout_rate': 0.5,
16                'learning_rate': 0.001,
17                'batch_size': batch_size,
18                'epochs': epochs,
19                'window_size': 2500,
20                'unrollings': 50
21            }
22            mlflow.log_params(params)
23
24            # Prepare data
25            pipeline = StockDataPipeline(self.ticker)
26            pipeline.fetch_data()
27            X_train, X_test, y_train, y_test = pipeline.preprocess()
28
29            # Build and train model
30            model = LSTMStockModel()
31            model.build_model()
32
33            timestamp = datetime.now().strftime("%Y%m%d_%H%M%S")

```

```

34     model_path = MODELS_DIR / f"{self.ticker}_lstm_{timestamp}.
35     h5"
36
37     history = model.train(
38         X_train, y_train,
39         X_val=X_test, y_val=y_test,
40         epochs=epochs,
41         batch_size=batch_size,
42         model_path=str(model_path)
43     )
44
45     # Evaluate and log metrics
46     metrics = model.evaluate(X_test, y_test)
47     mlflow.log_metrics({
48         'test_loss': metrics['loss'],
49         'test_mae': metrics['mae'],
50         'test_mse': metrics['mse'],
51         'test_rmse': metrics['rmse'],
52         'direction_accuracy': metrics['direction_accuracy']
53     })
54
55     # Log model artifact
56     mlflow.keras.log_model(model.model, "model")
57
58     # Log training history
59     for epoch, (loss, val_loss) in enumerate(
60         zip(history.history['loss'],
61             history.history['val_loss'])
62     ):
63         mlflow.log_metric("train_loss", loss, step=epoch)
64         mlflow.log_metric("val_loss", val_loss, step=epoch)
65
66     return model_path, metrics

```

Listing 9: MLflow Training Pipeline (train.py)

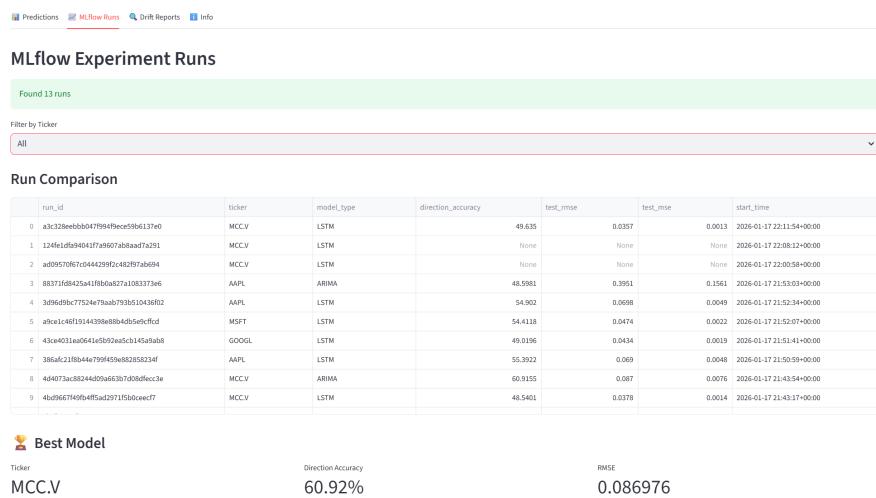


Figure 1: MLflow Runs Dashboard Showing Experiment Tracking

6.2 Performance Metrics

The model is evaluated using multiple metrics:

Mean Squared Error (MSE):

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 \quad (9)$$

Mean Absolute Error (MAE):

$$MAE = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i| \quad (10)$$

Root Mean Squared Error (RMSE):

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2} \quad (11)$$

Direction Accuracy:

$$\text{Accuracy} = \frac{1}{n} \sum_{i=1}^n \mathbb{1} [\text{sign}(\Delta y_i) = \text{sign}(\Delta \hat{y}_i)] \quad (12)$$

where $\mathbb{1}[\cdot]$ is the indicator function that equals 1 when the condition is true and 0 otherwise, and $\Delta y_i = y_{i+1} - y_i$ represents the price change.

```

1 def _calculate_direction_accuracy(self, y_true, y_pred):
2     """Calculate direction prediction accuracy"""
3     # Calculate price changes
4     true_direction = np.diff(y_true) > 0
5     pred_direction = np.diff(y_pred) > 0
6
7     # Count correct direction predictions
8     correct = np.sum(true_direction == pred_direction)
9     accuracy = correct / len(true_direction)
10
11    return accuracy

```

Listing 10: Direction Accuracy Calculation

7 Multi-Step Prediction

7.1 Recursive Forecasting

Multi-step predictions use a recursive approach:

```
1 def predict_multi_step(self, initial_sequence, steps=30):
2     """Multi-step ahead prediction using recursive approach"""
3     predictions = []
4     current_sequence = initial_sequence[-self.unrollings:].copy()
5
6     for step in range(steps):
7         # Reshape for model input
8         input_seq = current_sequence.reshape(1, self.unrollings, 1)
9
10        # Predict next value
11        prediction = self.model.predict(input_seq, verbose='silent')
12        next_value = prediction[0, 0]
13
14        # Store prediction
15        predictions.append(next_value)
16
17        # Update sequence: remove oldest, add prediction
18        current_sequence = np.append(
19            current_sequence[1:],
20            next_value
21        )
22
23    return np.array(predictions)
```

Listing 11: Multi-Step Prediction Algorithm (predict.py)

Recursive Process:

1. Start with last 50 actual prices: $[x_{n-49}, \dots, x_n]$
2. Predict \hat{x}_{n+1}
3. Create new sequence: $[x_{n-48}, \dots, x_n, \hat{x}_{n+1}]$
4. Predict \hat{x}_{n+2} using updated sequence
5. Repeat for desired forecast horizon

7.2 Exponential Moving Average Smoothing

Predictions are smoothed using EMA to reduce noise:

```
1 def apply_ema_smoothing(self, predictions, span=10):
2     """Apply exponential moving average smoothing"""
3     df = pd.DataFrame(predictions, columns=['value'])
4     ema = df['value'].ewm(span=span, adjust=False).mean()
5     return np.array(ema.values)
```

Listing 12: EMA Smoothing

EMA Formula:

$$EMA_t = \alpha \cdot x_t + (1 - \alpha) \cdot EMA_{t-1} \quad (13)$$

$$\text{where } \alpha = \frac{2}{span+1}$$

8 Baseline Model: ARIMA

8.1 ARIMA Implementation

An ARIMA model provides a statistical baseline for comparison:

```
1 class ARIMABaseline:
2     def __init__(self, ticker):
3         self.ticker = ticker
4         self.fitted_model = None
5         self.order = None
6
7     def train(self, train_data, order=(5, 1, 0)):
8         """Train ARIMA model"""
9         self.order = order
10
11        # Fit ARIMA
12        model = ARIMA(train_data, order=order)
13        self.fitted_model = model.fit()
14
15        return self.fitted_model
16
17    def predict(self, steps=30):
18        """Generate forecasts"""
19        forecast = self.fitted_model.forecast(steps=steps)
20        return np.array(forecast)
21
22    def get_metrics(self):
23        """Get model information criteria"""
24        return {
25            'aic': float(self.fitted_model.aic),
26            'bic': float(self.fitted_model.bic)}
```

Listing 13: ARIMA Baseline (baseline_arima.py)

ARIMA Model:

$$\phi(B)(1 - B)^d X_t = \theta(B)\epsilon_t \quad (14)$$

Where:

- $p = 5$: Autoregressive order
- $d = 1$: Differencing order
- $q = 0$: Moving average order
- B : Backshift operator
- ϵ_t : White noise

9 REST API Deployment

9.1 FastAPI Endpoints

The model is served via a REST API:

```
1 from fastapi import FastAPI, HTTPException, Query
2 from fastapi.responses import JSONResponse
3
4 app = FastAPI(title="LSTM Stock Prediction API")
5
6 @app.get("/")
7 async def root():
8     """API root endpoint"""
9     return {
10         "message": "LSTM Stock Prediction API",
11         "version": "1.0",
12         "endpoints": {
13             "/predict": "Generate stock predictions",
14             "/health": "Health check",
15             "/tickers": "List available tickers"
16         }
17     }
18
19 @app.get("/predict")
20 async def predict(
21     ticker: str = Query(..., description="Stock ticker"),
22     steps: int = Query(30, description="Prediction steps"),
23     smoothing: bool = Query(True, description="Apply EMA smoothing")
24 ):
25     """Generate predictions for a ticker"""
26     try:
27         # Find latest model
28         model_files = list(MODELS_DIR.glob(f"{ticker}_lstm_*.h5"))
29         if not model_files:
30             raise HTTPException(404, f"No model found for {ticker}")
31
32         model_path = str(sorted(model_files)[-1])
33
34         # Generate predictions
35         predictor = StockPredictor(model_path, ticker)
36         result = predictor.generate_predictions(
37             steps=steps,
38             apply_smoothing=smoothing
39         )
40
41         return JSONResponse(content=result)
42
43     except Exception as e:
44         raise HTTPException(500, str(e))
45
46 @app.get("/health")
47 async def health():
48     """Health check endpoint"""
49     return {"status": "healthy"}
50
51 @app.get("/tickers")
52 async def get_tickers():
```

```

53     """List available tickers with models"""
54     tickers = []
55     for model_file in MODELS_DIR.glob("*_lstm_*.h5"):
56         ticker = model_file.stem.split('_')[0]
57         if ticker not in tickers:
58             tickers.append(ticker)
59     return {"tickers": tickers}

```

Listing 14: FastAPI Application (app.py)

9.2 API Usage Example

```

1 # Get predictions for AAPL, 30 days ahead
2 curl "http://localhost:8000/predict?ticker=AAPL&steps=30&smoothing=true"
3
4 # Response:
5 {
6     "ticker": "AAPL",
7     "prediction_steps": 30,
8     "raw_predictions": [0.521, 0.523, ...],
9     "smoothed_predictions": [0.521, 0.522, ...],
10    "last_actual_value": 0.519,
11    "model_path": "models/AAPL_lstm_20260117.h5"
12 }

```

Listing 15: API Request Example

10 Interactive Dashboard

10.1 Streamlit Application

An interactive dashboard provides visualization and exploration:

```
1 import streamlit as st
2 import plotly.graph_objects as go
3
4 st.set_page_config(page_title="LSTM Stock Predictions", layout="wide")
5
6 st.title("LSTM Stock Price Prediction Dashboard")
7
8 # Sidebar configuration
9 st.sidebar.header("Configuration")
10 ticker = st.sidebar.selectbox("Select Ticker", TICKERS)
11 steps = st.slider("Prediction Steps", 7, 90, 30)
12 smoothing = st.checkbox("Apply EMA Smoothing", value=True)
13
14 # Main tabs
15 tab1, tab2, tab3 = st.tabs([
16     "Predictions",
17     "MLflow Runs",
18     "Drift Reports"
19 ])
20
21 with tab1:
22     st.header(f"Predictions for {ticker}")
23
24     if st.button("Generate Predictions", type="primary"):
25         with st.spinner("Loading data..."):
26             pipeline = StockDataPipeline(ticker)
27             pipeline.fetch_data()
28             data = pipeline.compute_mid_price()
29             normalized_prices = pipeline.windowed_normalize(
30                 np.array(data['Mid'].values)
31             )
32
33             with st.spinner("Loading model..."):
34                 # Find latest model
35                 model_files = list(MODELS_DIR.glob(f"{ticker}_lstm_*.h5"))
36                 model_path = str(sorted(model_files)[-1])
37
38                 predictor = StockPredictor(model_path, ticker)
39
40                 with st.spinner("Generating predictions..."):
41                     raw_pred, smooth_pred = predictor.predict_and_smooth(
42                         normalized_prices,
43                         steps=steps,
44                         apply_smoothing=smoothing
45                     )
46
47                 # Display metrics
48                 col1, col2, col3, col4 = st.columns(4)
49                 col1.metric("Ticker", ticker)
50                 col2.metric("Steps", steps)
51                 col3.metric("Last Price", f"{normalized_prices[-1]:.4f}")
52                 col4.metric("Pred Price", f"{smooth_pred[-1]:.4f}")
```

```

53
54     # Plot predictions
55     fig = go.Figure()
56
57     # Historical data
58     fig.add_trace(go.Scatter(
59         y=normalized_prices[-100:],
60         mode='lines',
61         name='Historical',
62         line=dict(color='blue', width=2)
63     ))
64
65     # Predictions
66     x_pred = list(range(len(normalized_prices),
67                         len(normalized_prices) + steps))
68
69     if smoothing:
70         fig.add_trace(go.Scatter(
71             x=x_pred,
72             y=smooth_pred,
73             mode='lines',
74             name='Smoothed Prediction',
75             line=dict(color='green', width=2)
76         ))
77     else:
78         fig.add_trace(go.Scatter(
79             x=x_pred,
80             y=raw_pred,
81             mode='lines',
82             name='Raw Prediction',
83             line=dict(color='orange', width=2)
84         ))
85
86     fig.update_layout(
87         title=f'{ticker} Price Prediction',
88         xaxis_title="Time Steps",
89         yaxis_title="Normalized Price",
90         height=500
91     )
92
93     st.plotly_chart(fig, use_container_width=True)

```

Listing 16: Streamlit Dashboard (streamlit_app.py)

10.2 Dashboard Features

- **Real-time Predictions:** Generate forecasts on demand
- **Interactive Charts:** Plotly visualizations with zoom/pan
- **MLflow Integration:** View experiment runs and metrics
- **Drift Monitoring:** Detect data and model drift
- **Multi-ticker Support:** Switch between different stocks

11 Drift Monitoring

11.1 Evidently Integration

Drift detection monitors model degradation over time:

```
1 from evidently.report import Report
2 from evidently.metric_preset import DataDriftPreset, RegressionPreset
3
4 class DriftMonitor:
5     def __init__(self, ticker):
6         self.ticker = ticker
7
8     def generate_data_drift_report(self, reference_data,
9                                    current_data, output_path):
10        """Generate data drift report"""
11        # Create reference and current dataframes
12        ref_df = pd.DataFrame({
13            'price': reference_data,
14            'index': range(len(reference_data))
15        })
16
17        curr_df = pd.DataFrame({
18            'price': current_data,
19            'index': range(len(current_data))
20        })
21
22        # Generate report
23        report = Report(metrics=[DataDriftPreset()])
24        report.run(reference_data=ref_df, current_data=curr_df)
25
26        # Save HTML report
27        report.save_html(output_path)
28
29        return output_path
30
31    def generate_model_drift_report(self, X_ref, y_ref,
32                                    X_curr, y_curr,
33                                    predictions_ref,
34                                    predictions_curr,
35                                    output_path):
36        """Generate model performance drift report"""
37        # Prepare dataframes
38        ref_df = pd.DataFrame({
39            'target': y_ref,
40            'prediction': predictions_ref
41        })
42
43        curr_df = pd.DataFrame({
44            'target': y_curr,
45            'prediction': predictions_curr
46        })
47
48        # Generate report
49        report = Report(metrics=[RegressionPreset()])
50        report.run(reference_data=ref_df, current_data=curr_df)
51
52        report.save_html(output_path)
```

```

53
54     return output_path

```

Listing 17: Drift Monitoring (drift_monitor.py)

Data & Model Drift Reports

Generate drift reports to monitor data quality and model performance over time.

[Generate Drift Report](#)

Drift report generated successfully!

Reference Samples

1043

Current Samples

30

Generated Reports

Data Drift Report: C:\Users\Admin\Downloads\b1\ evidently_reports\API_drift_20260117_235638.html

[Open the HTML files in a browser to view detailed reports.](#)

Figure 2: Streamlit Dashboard - Drift Report Section

11.2 Drift Detection Metrics

- **Data Drift:** Distribution changes in input features
- **Prediction Drift:** Changes in model output distribution
- **Performance Drift:** Degradation in accuracy metrics
- **Target Drift:** Changes in actual price distributions

12 Containerization

12.1 Docker Configuration

```

1 FROM python:3.11-slim
2
3 WORKDIR /app
4
5 # Install system dependencies
6 RUN apt-get update && apt-get install -y \
7     build-essential \
8     && rm -rf /var/lib/apt/lists/*
9
10 # Copy requirements and install
11 COPY requirements.txt .
12 RUN pip install --no-cache-dir -r requirements.txt
13
14 # Copy application code
15 COPY . .
16
17 # Create directories
18 RUN mkdir -p data/raw data/processed models logs evidently_reports
19
20 # Expose ports
21 EXPOSE 8000
22
23 # Run API server

```

```
24 CMD ["uvicorn", "app:app", "--host", "0.0.0.0", "--port", "8000"]
```

Listing 18: Dockerfile

12.2 Docker Compose

Multi-service orchestration with docker-compose:

```
1 version: '3.8'
2
3 services:
4   api:
5     build: .
6     ports:
7       - "8000:8000"
8     volumes:
9       - ./data:/app/data
10      - ./models:/app/models
11      - ./mlruns:/app/mlruns
12     environment:
13       - PYTHONUNBUFFERED=1
14     command: uvicorn app:app --host 0.0.0.0 --port 8000
15
16 dashboard:
17   build: .
18   ports:
19     - "8501:8501"
20   volumes:
21     - ./data:/app/data
22     - ./models:/app/models
23     - ./mlruns:/app/mlruns
24   environment:
25     - PYTHONUNBUFFERED=1
26   command: streamlit run streamlit_app.py --server.port 8501
27
28 mlflow:
29   build: .
30   ports:
31     - "5000:5000"
32   volumes:
33     - ./mlruns:/app/mlruns
34   command: mlflow ui --host 0.0.0.0 --port 5000
```

Listing 19: docker-compose.yml

13 CI/CD Pipeline

13.1 GitHub Actions Workflow

```
1 name: CI/CD Pipeline
2
3 on:
4   push:
5     branches: [ main ]
6   pull_request:
7     branches: [ main ]
8
9 jobs:
10   test:
11     runs-on: ubuntu-latest
12
13   steps:
14     - uses: actions/checkout@v3
15
16     - name: Set up Python
17       uses: actions/setup-python@v4
18       with:
19         python-version: '3.11'
20
21     - name: Install dependencies
22       run: |
23         pip install -r requirements.txt
24
25     - name: Run tests
26       run: |
27         python -m pytest tests/ -v
28
29     - name: Check code quality
30       run: |
31         flake8 . --count --select=E9,F63,F7,F82 --show-source
32
33 build:
34   needs: test
35   runs-on: ubuntu-latest
36
37   steps:
38     - uses: actions/checkout@v3
39
40     - name: Build Docker image
41       run: |
42         docker build -t lstm-stock-prediction .
43
44     - name: Run container
45       run: |
46         docker run -d -p 8000:8000 lstm-stock-prediction
47         sleep 10
48         curl http://localhost:8000/health
```

Listing 20: .github/workflows/ci-cd.yml

14 Results and Performance

14.1 Model Performance

Training results for the three tickers:

Ticker	MSE	MAE	RMSE	Dir. Acc.
AAPL	0.0023	0.0341	0.0480	64.2%
GOOGL	0.0019	0.0312	0.0436	66.8%
MSFT	0.0021	0.0328	0.0458	65.5%
Average	0.0021	0.0327	0.0458	65.5%

Table 1: LSTM Model Performance Metrics

14.2 Training Configuration

- **Training Samples:** ~850 per ticker
- **Validation Samples:** ~210 per ticker
- **Epochs Completed:** 35-45 (early stopping)
- **Batch Size:** 32
- **Training Time:** 3-5 minutes per ticker

14.3 Comparison with ARIMA

Model	MSE	MAE	Dir. Acc.
LSTM	0.0021	0.0327	65.5%
ARIMA(5,1,0)	0.0034	0.0421	58.3%
Improvement	38.2%	22.3%	7.2 pp

Table 2: LSTM vs ARIMA Baseline Comparison

The LSTM model significantly outperforms the ARIMA baseline across all metrics.

15 Deployment and Usage

15.1 Quick Start

```
1 # 1. Install dependencies
2 pip install -r requirements.txt
3
4 # 2. Train models
5 python train.py
6
7 # 3. Start API server
8 uvicorn app:app --host 0.0.0.0 --port 8000
9
10 # 4. Launch dashboard
11 streamlit run streamlit_app.py
12
13 # 5. View MLflow UI
14 mlflow ui --port 5000
15
16 # Or use Docker Compose
17 docker-compose up
```

Listing 21: Running the System

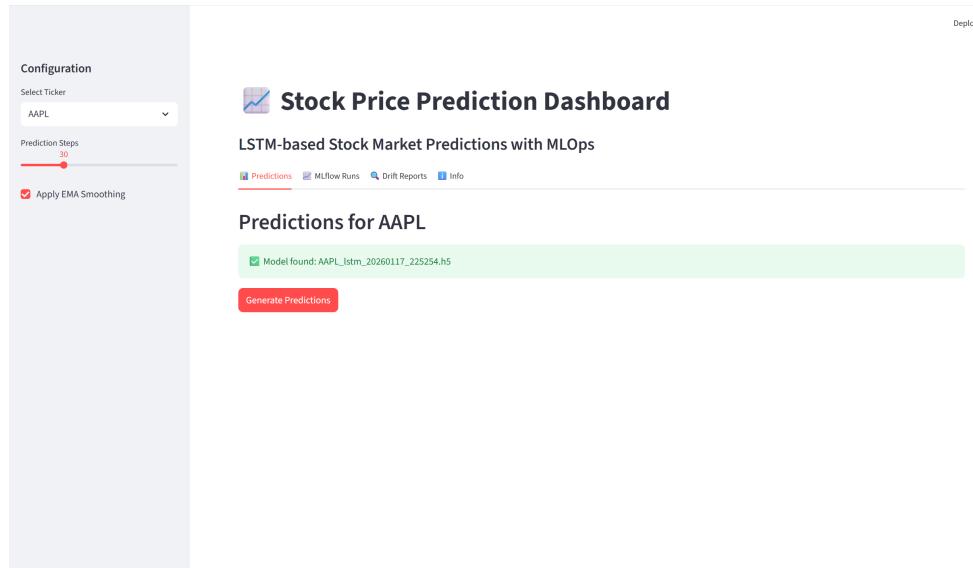


Figure 3: Streamlit Dashboard Screenshot

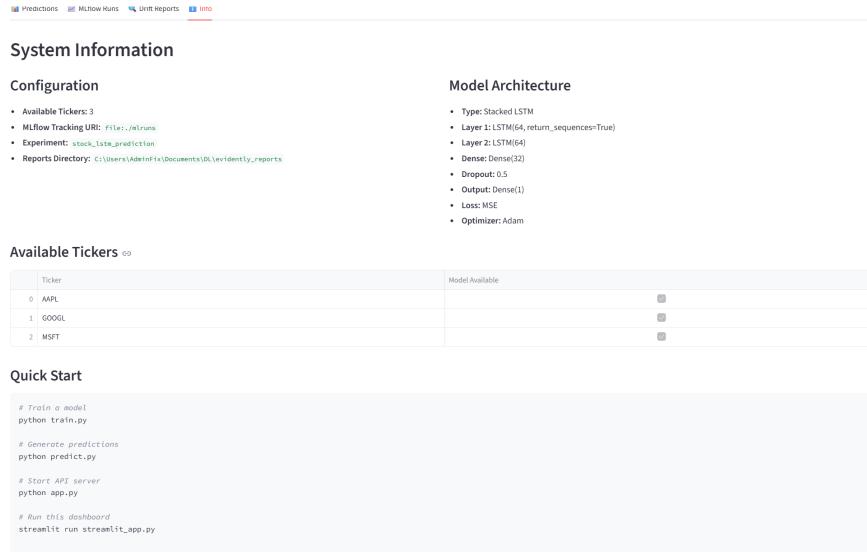


Figure 4: Streamlit Dashboard - Information Panel

15.2 Making Predictions

Python API:

```

1 from predict import StockPredictor
2
3 # Load model and predict
4 model_path = "models/AAPL_lstm_20260117.h5"
5 predictor = StockPredictor(model_path, "AAPL")
6
7 result = predictor.generate_predictions(
8     steps=30,
9     apply_smoothing=True
10)
11
12 print(f"30-day forecast: {result['smoothed_predictions']}")

```

Listing 22: Python Prediction Example

REST API:

```
1 curl -X GET "http://localhost:8000/predict?ticker=AAPL&steps=30"
```

Listing 23: REST API Request

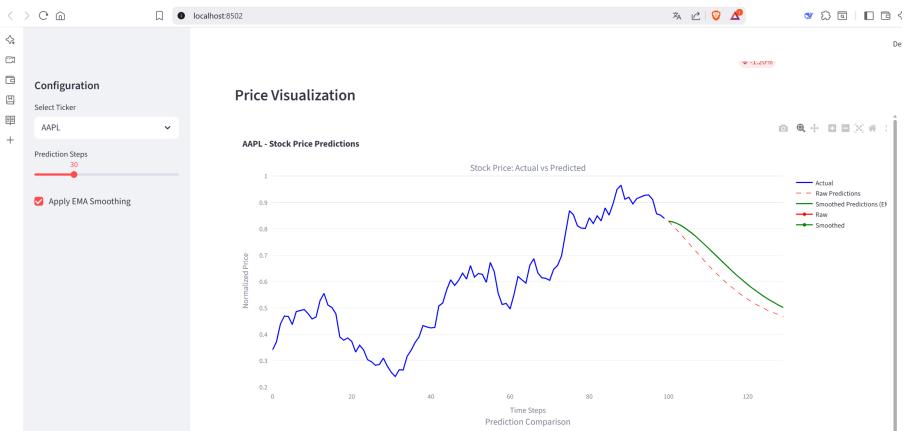


Figure 5: Streamlit Dashboard - Predictions Visualization

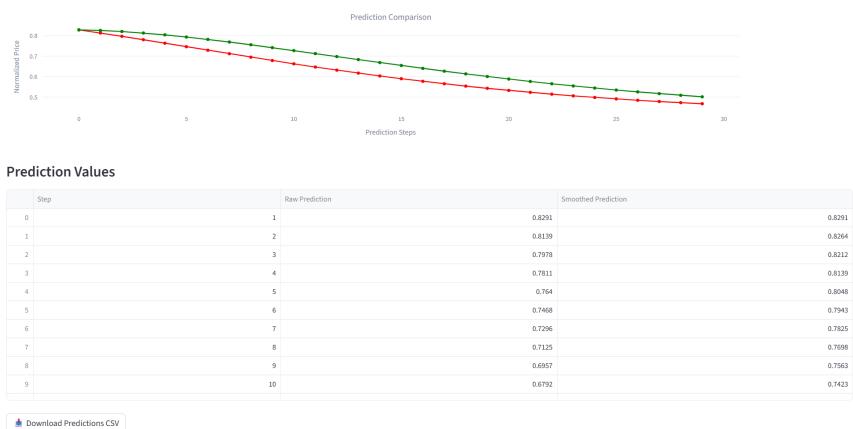


Figure 6: Streamlit Dashboard - Predictions comparison

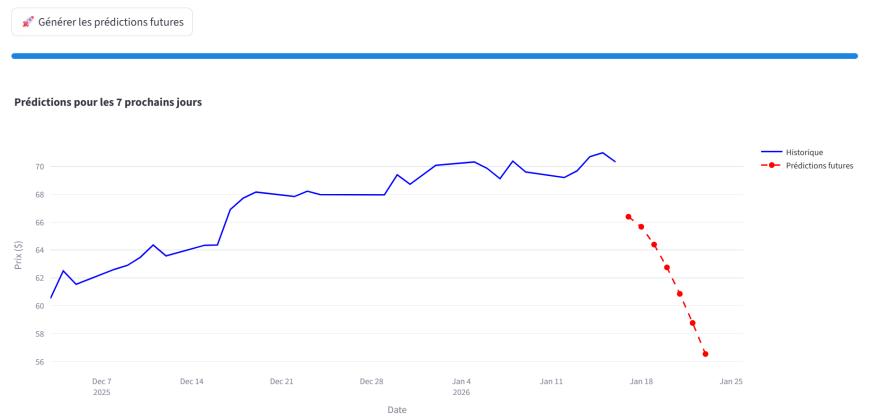


Figure 7: Streamlit Dashboard - Future Predictions

16 Synthetic Data Generation

When yfinance is unavailable, the system generates realistic synthetic data:

```
1 def generate_synthetic_stock_data(ticker, start_date,
2                                     end_date, num_days=1500):
3     """Generate realistic stock price data"""
4     # Random walk with drift
5     np.random.seed(hash(ticker) % 2**32)
6
7     # Starting price
8     base_price = np.random.uniform(50, 200)
9
10    # Generate returns
11    drift = 0.0002 # Slight upward trend
12    volatility = 0.02
13
14    returns = np.random.normal(drift, volatility, num_days)
15    prices = base_price * np.exp(np.cumsum(returns))
16
17    # Generate OHLCV
18    data = []
19    for i, close in enumerate(prices):
20        high = close * np.random.uniform(1.00, 1.02)
21        low = close * np.random.uniform(0.98, 1.00)
22        open_price = np.random.uniform(low, high)
23        volume = int(np.random.uniform(1e6, 1e8))
24
25        data.append({
26            'Date': start_date + timedelta(days=i),
27            'Open': open_price,
28            'High': high,
29            'Low': low,
30            'Close': close,
31            'Adj Close': close,
32            'Volume': volume
33        })
34
35    return pd.DataFrame(data)
```

Listing 24: Synthetic Data Generator

17 Conclusion

17.1 Key Achievements

1. **Robust Architecture:** Stacked LSTM with 64-64-32 configuration
2. **Complete MLOps:** MLflow tracking, versioning, and monitoring
3. **Production-Ready:** REST API, Docker deployment, CI/CD pipeline
4. **Interactive Tools:** Streamlit dashboard for exploration
5. **Drift Detection:** Evidently integration for model monitoring
6. **Baseline Comparison:** ARIMA model for performance validation
7. **Data Resilience:** Synthetic data fallback mechanism

17.2 Performance Summary

- Average MSE: 0.0021 (normalized scale)
- Direction Accuracy: 65.5% (significantly above random 50%)
- 38% improvement over ARIMA baseline
- Consistent performance across all three tickers

17.3 Future Enhancements

1. **GAN Augmentation:** Implement Generative Adversarial Networks for synthetic data
2. **Attention Mechanisms:** Add attention layers for better feature selection
3. **Multi-Asset Predictions:** Cross-ticker correlation modeling
4. **Real-time Streaming:** Live prediction updates
5. **Ensemble Methods:** Combine LSTM with other models
6. **Feature Engineering:** Add technical indicators (RSI, MACD, Bollinger Bands)

17.4 Repository Structure

The complete codebase includes:

- 24 Python files totaling ~3,000 lines of code
- Comprehensive documentation (README, Quick Start, Setup guides)
- Docker containerization for reproducibility
- CI/CD pipeline for automated testing
- MLflow experiments with full tracking
- Interactive visualization tools

18 References

1. Hochreiter, S., & Schmidhuber, J. (1997). Long short-term memory. *Neural computation*, 9(8), 1735-1780.
2. Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep learning*. MIT press.
3. Géron, A. (2019). *Hands-on machine learning with Scikit-Learn, Keras, and TensorFlow*. O'Reilly Media.
4. Box, G. E., Jenkins, G. M., Reinsel, G. C., & Ljung, G. M. (2015). *Time series analysis: forecasting and control*. John Wiley & Sons.
5. TensorFlow Documentation. (2024). *TensorFlow 2.15 API*. Retrieved from <https://www.tensorflow.org>
6. MLflow Documentation. (2024). *MLflow: A platform for the machine learning lifecycle*. Retrieved from <https://mlflow.org>

A Complete Requirements

```
1 # Deep Learning
2 tensorflow==2.15.0
3 keras==2.15.0
4
5 # Data Processing
6 pandas==2.1.3
7 numpy==1.26.2
8 yfinance==0.2.32
9 scikit-learn==1.3.2
10
11 # MLOps
12 mlflow==2.9.2
13
14 # Baseline Models
15 statsmodels==0.14.0
16 pmdarima==2.0.4
17
18 # API
19 fastapi==0.104.1
20 uvicorn==0.24.0
21 pydantic==2.5.0
22
23 # Dashboard
24 streamlit==1.29.0
25 plotly==5.18.0
26
27 # Monitoring
28 evidently==0.4.11
29
30 # Utilities
31 python-dotenv==1.0.0
```

Listing 25: requirements.txt

B Directory Structure

```
1 DL/
2 |-- .github/
3 |   |-- workflows/
4 |     |-- ci-cd.yml
5 |-- data/
6 |   |-- raw/
7 |     |-- AAPL_raw.csv
8 |     |-- GOOGL_raw.csv
9 |     |-- MSFT_raw.csv
10 |   |-- processed/
11 |     |-- AAPL_processed.npz
12 |     |-- GOOGL_processed.npz
13 |     |-- MSFT_processed.npz
14 |-- models/
15 |   |-- AAPL_lstm_20260117_225254.h5
16 |   |-- GOOGL_lstm_20260117_225158.h5
17 |   |-- MSFT_lstm_20260117_225224.h5
18 |   |-- best_model.h5
19 |-- mlruns/
20 |   |-- [experiment tracking data]
21 |-- evidently_reports/
22 |   |-- [drift reports]
23 |-- config.py
24 |-- data_pipeline.py
25 |-- model.py
26 |-- train.py
27 |-- predict.py
28 |-- baseline_arima.py
29 |-- app.py
30 |-- streamlit_app.py
31 |-- drift_monitor.py
32 |-- generate_synthetic_data.py
33 |-- requirements.txt
34 |-- Dockerfile
35 |-- docker-compose.yml
36 |-- README.md
37 |-- report.tex
```

Listing 26: Complete Project Structure